

MLP inference latency across kernels

